

Types of Learning

- Supervised
 - *Structured/labeled data*
 - Ex: picture \rightarrow label
 - Ex: picture \rightarrow digit number
- Unsupervised
 - *Unstructured data*
 - Dataset of pictures(not labeled)
- m denotes # of training examples

Binary Classification

- Classifies whether the input (is/has) or (isn't/doesn't have) a particular thing.
- Ex: Given an image, tell whether it has a cat in it or not 1(cat) vs 0(non-cat).
- In general, computers store images in 3 matrices(red, green, and blue) each of size width x height
- Images are stretched into a single column vector, usually starting with all the red pixels then all the green ones, then all the blue ones.
- n_x = the dimension of the input features
 - Ex: $64 \times 64 \times 3 = n_x = 12288$
- A single training example is denoted as:
 - (x, y) where $x \in \mathbb{R}^{n_x}$ and $y \in \{0, 1\}$
- m training examples: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
 - Other common notation:
 - * $m = m_{train}$
 - * $m_{test} = \#$ test examples
- X is used to denote a matrix of all the inputs
 - $X \in \mathbb{R}^{n_x \times m}$
 - $X.shape = (n_x, m)$
- Y is used to denote a row vector of all the labels
 - $Y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]$
 - $Y \in \mathbb{R}^{1 \times m}$
 - $Y.shape = (1, m)$

Logistic Regression

- Given x , we want $\hat{y} = P(y = 1 | x)$
- Parameters:
 - $x \in \mathbb{R}^{n_x}$
 - $b \in \mathbb{R}$
- Output:
 - $\sigma(w^T x + b)$

- We use the *sigmoid* function to make sure that the output is between 0 and 1, and centered on 0.5

Logistic Regression Cost Function

- Given $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$, we want $\hat{y}^{(i)} \approx y^{(i)}$
- Loss(error) function:
 - Computes the error for a single training example
 - $L(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log (1 - \hat{y}))$
 - If $y = 1$: $\mathcal{L}(\hat{y}, y) = -\log \hat{y} \leftarrow$ wants \hat{y} to be large
 - If $y = 0$: $\mathcal{L}(\hat{y}, y) = -\log (1 - \hat{y}) \leftarrow$ wants \hat{y} to be small
- Cost function:
 - Averages the loss of the entire training set
 - $J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$

Gradient Descent

- $w := w - \alpha \frac{dJ(w, b)}{dw}$
- $b := b - \alpha \frac{dJ(w, b)}{db}$
- α is called the **Learning Rate**

Neural Network Representation

- Each layer is denoted by $a^{[n]}$
 - a is a column vector representing activations of neurons
 - The superscript n represents the layer that we are looking at
- The input layer is denoted $a^{[0]}$

Activation Functions

- Activation functions are denoted $g^{[n]}$
 - g represents the function
 - n represents the layer
- Previously, we used Sigmoid(σ)
 - Output is always between 0 and 1
 - Centered on 0.5
- Tanh is another alternative
 - Output is always between -1 and 1
 - Centered on 0
- Tanh is usually much better except for the output layer for binary activation
- ReLU = $\max(0, z)$
 - Rectified Linear Unit
- Leaky ReLU is another option
- No activation function is called a **linear Activation** function.
 - Its very uncommon to use a linear activation function.

- Sometimes used for the output layer when $y \in \mathbb{R}$

L-layer Deep Neural Network

- Forward Pass
 - $Z^{[\ell]} = W^{[\ell]} \cdot A^{[\ell-1]} + b^{[\ell]}$
 - $A^{[\ell]} = g^{[\ell]}(Z^{[\ell]})$
- Backward Pass
 - $dZ^{[\ell]} = dA^{[\ell]} * g^{[\ell]'}(Z^{[\ell]})$
 - $dW^{[\ell]} = \frac{1}{m} dZ^{[\ell]} \cdot A^{[\ell-1]T}$
 - $db^{[\ell]} = \frac{1}{m} \text{np.sum}(dZ^{[\ell]}, \text{axis}=1)$
 - $dA^{[\ell-1]} = W^{[\ell]T} \cdot dZ^{[\ell]}$

Train/Dev/Test Sets

- For small datasets, you have to allocate a larger % of examples for testing and validation.
 - Ex: 70/30 train/test split
- For larger datasets, a much larger % of examples are used for testing.
 - Ex: 99/1/1 train/dev/test split
- Always make sure the dev and test sets come from the same distribution.

Bias/Variance

- High **Bias** means *underfitting*.
 - Ex: 15% training error and 16% dev error
- High **Variance** means *overfitting*.
 - Ex: 1% training error and 11% dev error.
- High bias and high variance can both be present.
 - Ex: 15% training error and 30 % dev error.
- In between high bias and high variance is **just right**
 - Ex: 0.5% training error and 1% dev error.

Basic Recipe for Machine Learning

- Issues with high bias?
 - Increase the size of your network
 - Train longer
- High variance?
 - Try to get more data
 - Regularization
 - More appropriate architecture

L2 Regularization

- **L2 Regularization**
 - Sum of all the weights squared

$$\|w\|_2^2 = \sum_{j=1}^{n_x} w_j^2 = w^T w$$

Dropout Regularization

- keep_prob = 0.8
- d3 = np.random.randn(a3.shape[0], a3.shape[1]) < keep_prob
- a3 = np.multiply(a3, d3)
- a3 /= keep_prob
 - This is called *Inverted Dropout*
 - This line ensures that the values are bumped up for the next layer

Normalizing Inputs

- Normalizing Inputs allows us to use larger learning rates because we won't be oscillating back and forth in our loss function
- Subtract Mean
 - $\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$
 - $x := x - \mu$
- Normalize Variance
 - $\sigma^2 = \frac{1}{m} \sum_{i=1}^m x * x$
 - * element-wise squaring
 - $x /= \sigma$

Gradient Check

- Compute the gradient, d/θ , then compute the limit of the gradient and compare the two in euclidean distance
- 10^{-7} is usually considered a fine margin of difference
- 10^{-5} is questionable but may be fine
- 10^{-3} is very likely that something is wrong

Mini-batch Gradient Descent

- t denotes the mini-batch index we are using for our dataset
- $t = 1, 2, \dots, 5000$ for a dataset split into 5000 batches
- $Z^{[1]} = W^{[1]}X^{(t)} + b^{[1]}$
- $A^{[1]} = g^{[1]}(Z^{[1]})$
- ... and so on

- Cost is then computed by dividing by the size of the batch instead of $\frac{1}{m}$
- *epoch* is defined as a single pass through the training set
- Tends to make the cost-graph oscillate a bit
- Batch size is a hyper-parameter between 64 and 512
 - Usually a power of 2

Batch Gradient Descent

- Mini-batch Gradient Descent with a batch size of m
- Very large steps
- Takes too long per iteration

Stochastic Gradient Descent(SGD) - Mini-batch Gradient Descent with a batch size of 1 - Tends to be very noisy and wanders quite a bit - Never exactly converges - Lose speedup from vectorization**

Exponentially Weighted Average

- $V_t = \beta V_{t-1} + (1 - \beta)\theta_t$
- Averages your data over the last $\frac{1}{1-\beta}$ datapoints
- β is usually between 0 and 1
- Smooths out a pretty noisy graph(like a graph of temperature)
- The higher the value of β , the slower it is to react, so keep that in mind
- the lower the value of β , the faster it is to react, but the more noisy it gets

Gradient Descent with Momentum

- $V_{dW} = \beta V_{dW} + (1 - \beta)dW$
- $V_{db} = \beta V_{db} + (1 - \beta)db$
- $W = W - \alpha V_{dW}$
- $b = b - \alpha V_{db}$
- You can use momentum to make bigger weight updates at first then smaller updates later on

Learning Rate Decay

- Decreasing the learning rate as time goes on, so fewer and fewer steps are being made
- Helps convergence at the end while having big steps at the beginning
- $\alpha = \frac{1}{1+\text{decay-rate}} * \text{epoch-num}\alpha_0$
- The learning rate becomes a function of the epoch number

Hyperparameter Tuning

- Start with learning rate α
- Then try momentum β
- Then try # hidden units
- Then try mini batch size
- Lastly, try # of layers and the learning rate decay.
- Use random values. Do not use a *grid*
- Use logarithmic scale to make sure you are focusing resources where they are most likely best
- `$r = -4 * $ np.random.rand()` which makes $r \in [-4, 0]$
- $\alpha = 10^r$
- Try to train multiple models at once and cull them as they become clearly bad
 - The *caviar* approach
- Train one model at a time and fidget with settings till you get it right
 - The *panda* approach